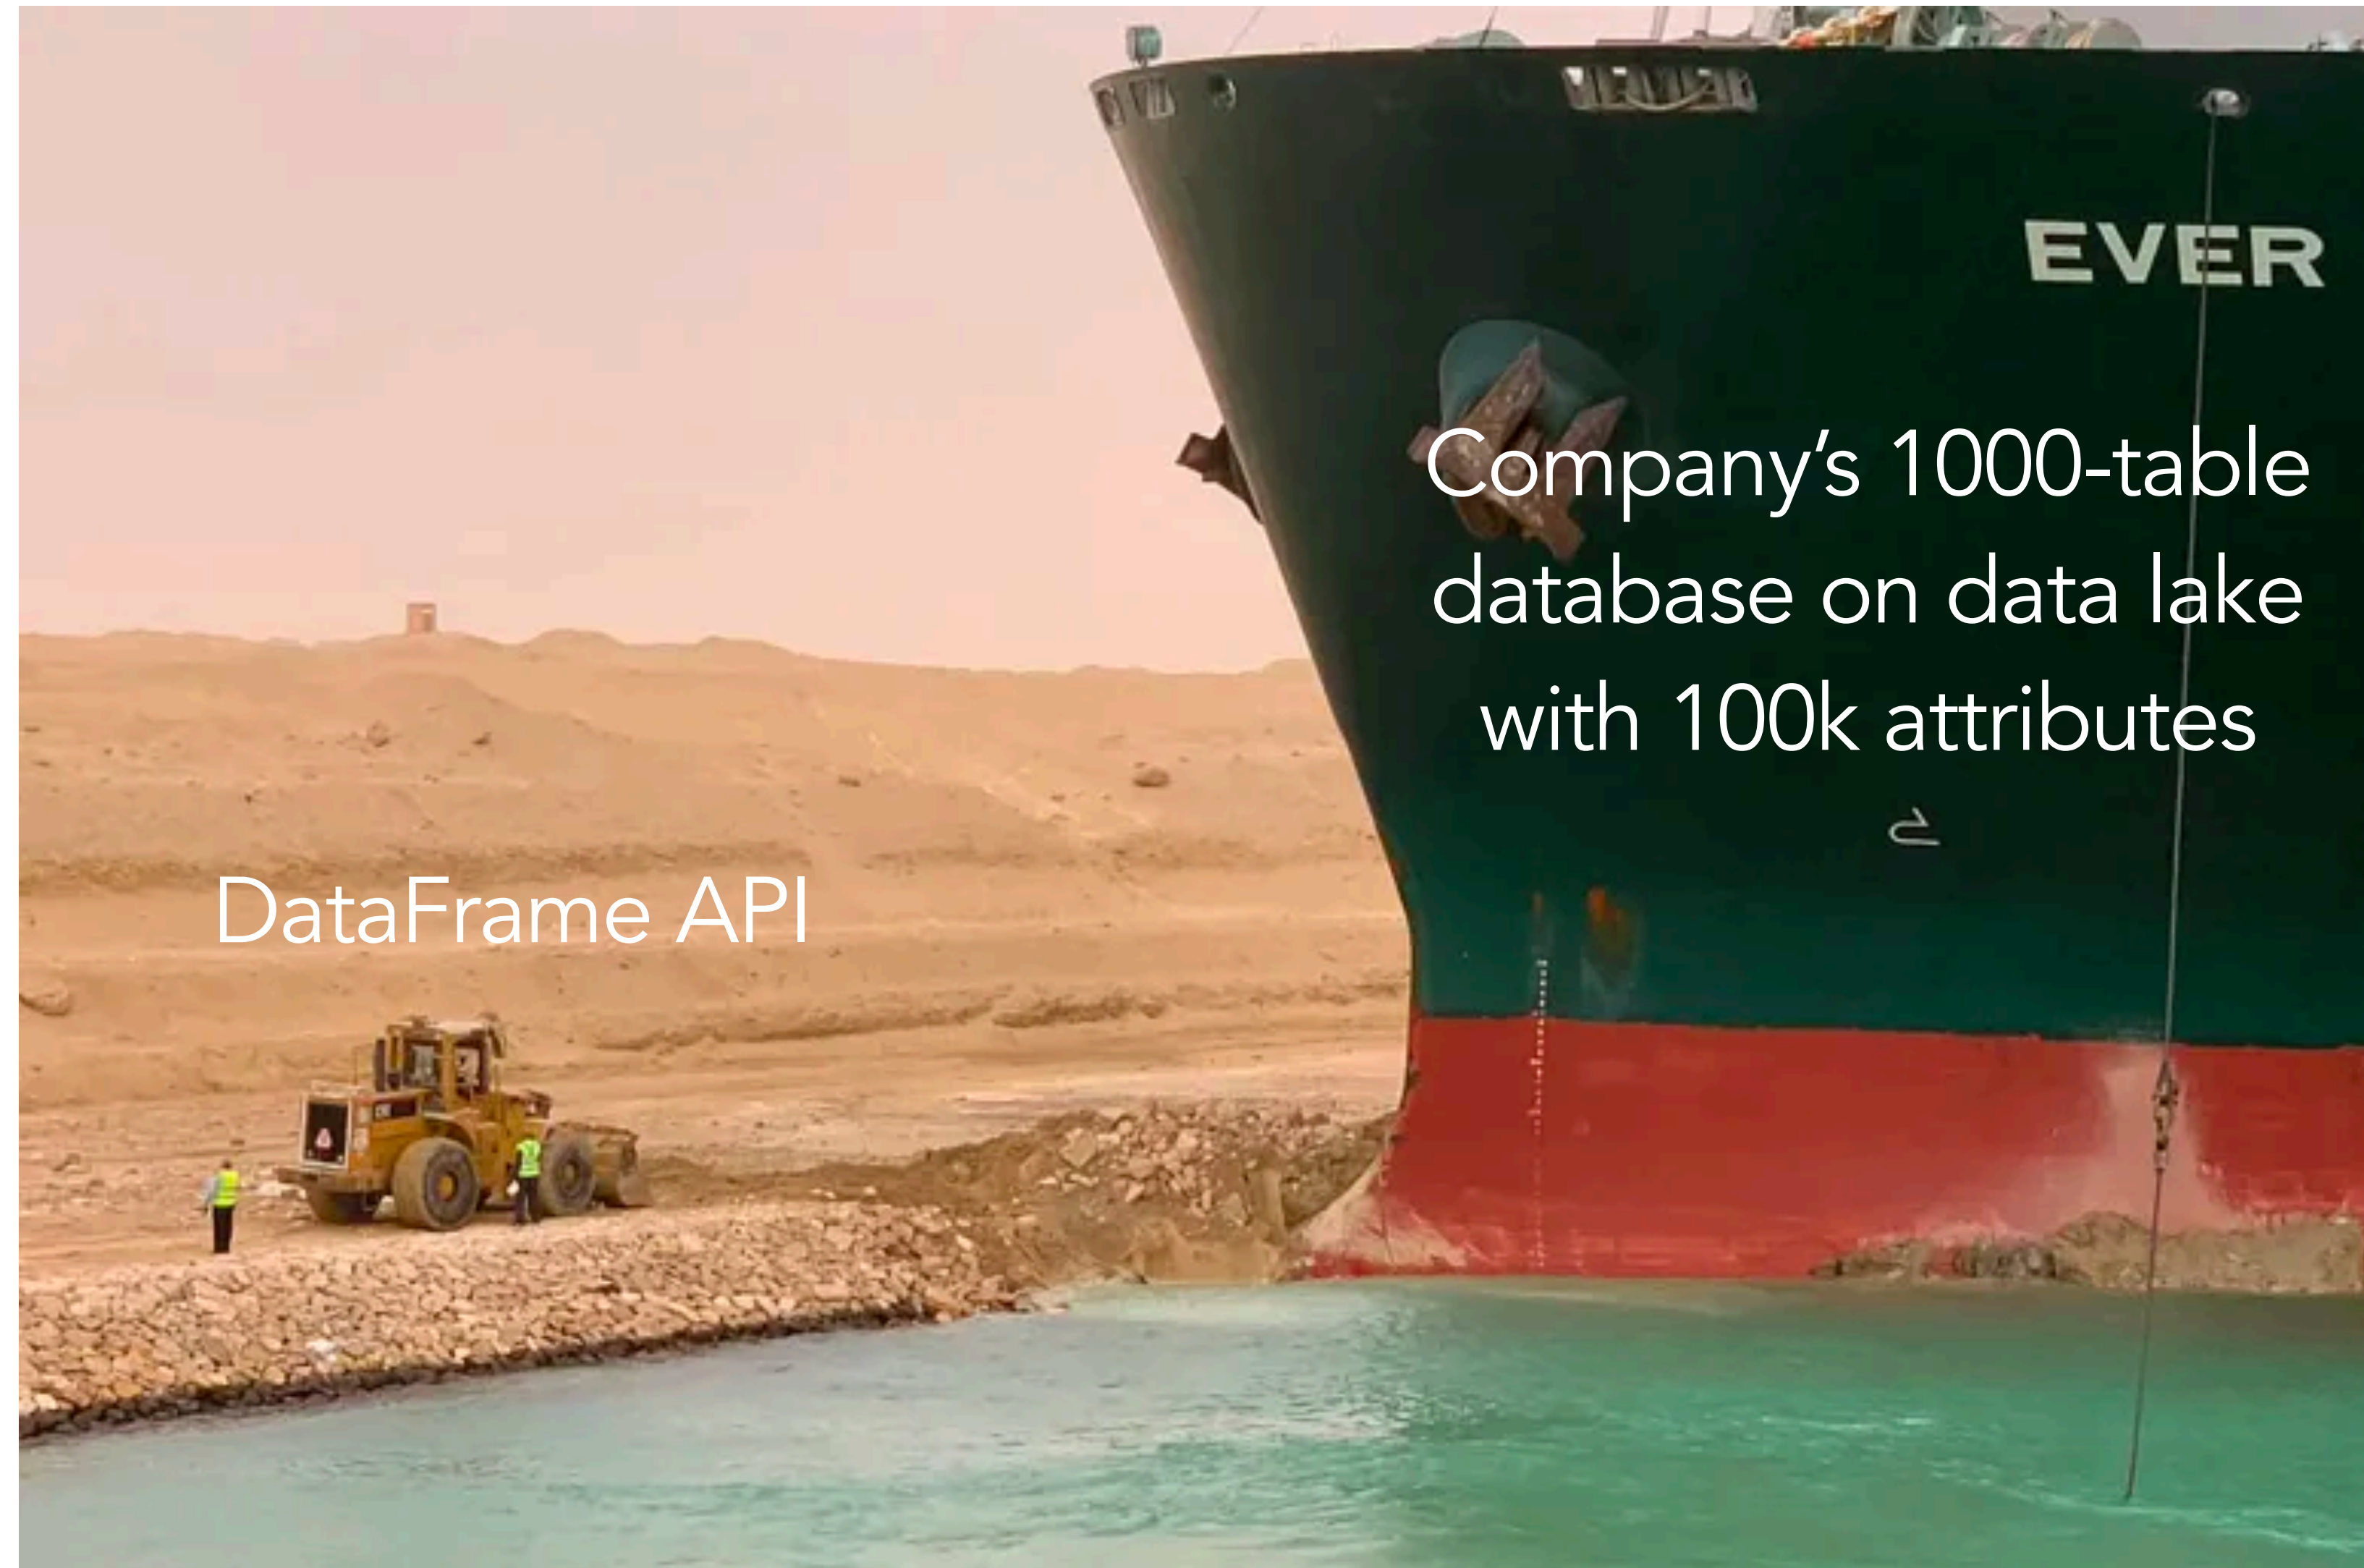


DSC 204a Scalable Data Systems

- Haojian Jin



Company's 1000-table
database on data lake
with 100k attributes

DataFrame API

Where are we in the class?

Foundations of Data Systems (2 weeks)

- Digital representation of Data → Computer Organization → Memory hierarchy → Process → Storage

Scaling Distributed Systems (3 weeks)

- Cloud → Network → Distributed storage → Parallelism → Partition and replication

Data Processing and Programming model (5 weeks)

- Data Models evolution → **Data encoding evolution** → IO & Unix Pipes → Batch processing (MapReduce) → Stream processing (Spark)

Today's topic: Data Encoding

- **Formats for Encoding Data**
 - Language-Specific Formats
 - JSON, XML, and Binary Variants
 - BINARY ENCODING
- Modes of dataflow
 - Database
 - REST
 - RPC
 - GraphQL
- Summary

Why encoding?

- Data in memory
 - e.g., objects, structs, lists, arrays, hash tables, trees
 - Efficient access and manipulation by the CPU (typically using pointers)
 - Why pointers? => Random address access
- Data in storage or network
 - No pointers.
 - Self-contained sequence of bytes.


```
num_tests = 10

obj = np.random.normal(0.5, 1, [240, 320, 3])

command = 'pickle.dumps(obj)'
setup = 'from __main__ import pickle, obj'
result = timeit.timeit(command, setup=setup, number=num_tests)
print("pickle: %f seconds" % result)
```

```
pickle          : 0.847938 seconds
cPickle         : 0.810384 seconds
cPickle highest: 0.004283 seconds
json            : 1.769215 seconds
msgpack         : 0.270886 seconds
```

Language-Specific Formats

- Java: `java.io.Serializable`;
- Python has pickle;
- Pros:
 - Convenient: in-memory **objects** to be saved and restored.
- Cons:
 - Tied to a programming language.
 - Decoding may lead to over-privileged behaviors.
 - e.g., remote execution.
 - Versioning, forward and backward compatibility
 - Efficiency
- Summary: quick, dirty, small individual projects

JSON, XML, CSV

- Python: Json dump.
- JSON, XML, CSV: human-readable but verbose.
- JSON: browser friendly and simple
- Common cons:
 - too verbose and unnecessarily complicated
 - ambiguity around the encoding of numbers
 - XML and CSV don't distinguish a number and a string that happens to consist of digits
 - JSON doesn't distinguish integers and floating-point numbers, and it doesn't specify a precision.

JSON, XML, CSV

- Python: Json dump.
- Common cons:
 - JSON and XML have good support for Unicode character strings.
 - There is optional schema support for both XML and JSON
 - CSV does not have any schema,

Example

```
{  
  "userName": "Martin",  
  "favoriteNumber": 1337,  
  "interests": ["daydreaming", "hacking"]  
}
```

MessagePack

```
{  
  "userName": "Martin",  
  "favoriteNumber": 1337,  
  "interests": ["daydreaming", "hacking"]  
}
```

MessagePack

Byte sequence (66 bytes):

83	a8	75	73	65	72	4e	61	6d	65	a6	4d	61	72	74	69	6e	ae	66	61
76	6f	72	69	74	65	4e	75	6d	62	65	72	cd	05	39	a9	69	6e	74	65
72	65	73	74	73	92	ab	64	61	79	64	72	65	61	6d	69	6e	67	a7	68
61	63	6b	69	6e	67														

Breakdown:

object (3 entries)	string (length 8)	u	s	e	r	N	a	m	e	string (length 6)	M	a	r	t	i	n			
83	a8	75	73	65	72	4e	61	6d	65	a6	4d	61	72	74	69	6e			
	string (length 14)	f	a	v	o	r	i	t	e	N	u	m	b	e	r				
	ae	66	61	76	6f	72	69	74	65	4e	75	6d	62	65	72				
	uint16 1337	05	39							string (length 9)	i	n	t	e	r	e	s	t	s
	cd									a9	69	6e	74	65	72	65	73	74	73
array (2 entries)	string (length 11)	d	a	y	d	r	e	a	m	i	n	g							
92	ab	64	61	79	64	72	65	61	6d	69	6e	67							
	string (length 7)	h	a	c	k	i	n	g											
	a7	68	61	63	6b	69	6e	67											

Thrift BinaryProtocol

```
{  
  "userName": "Martin",  
  "favoriteNumber": 1337,  
  "interests": ["daydreaming", "hacking"]  
}
```

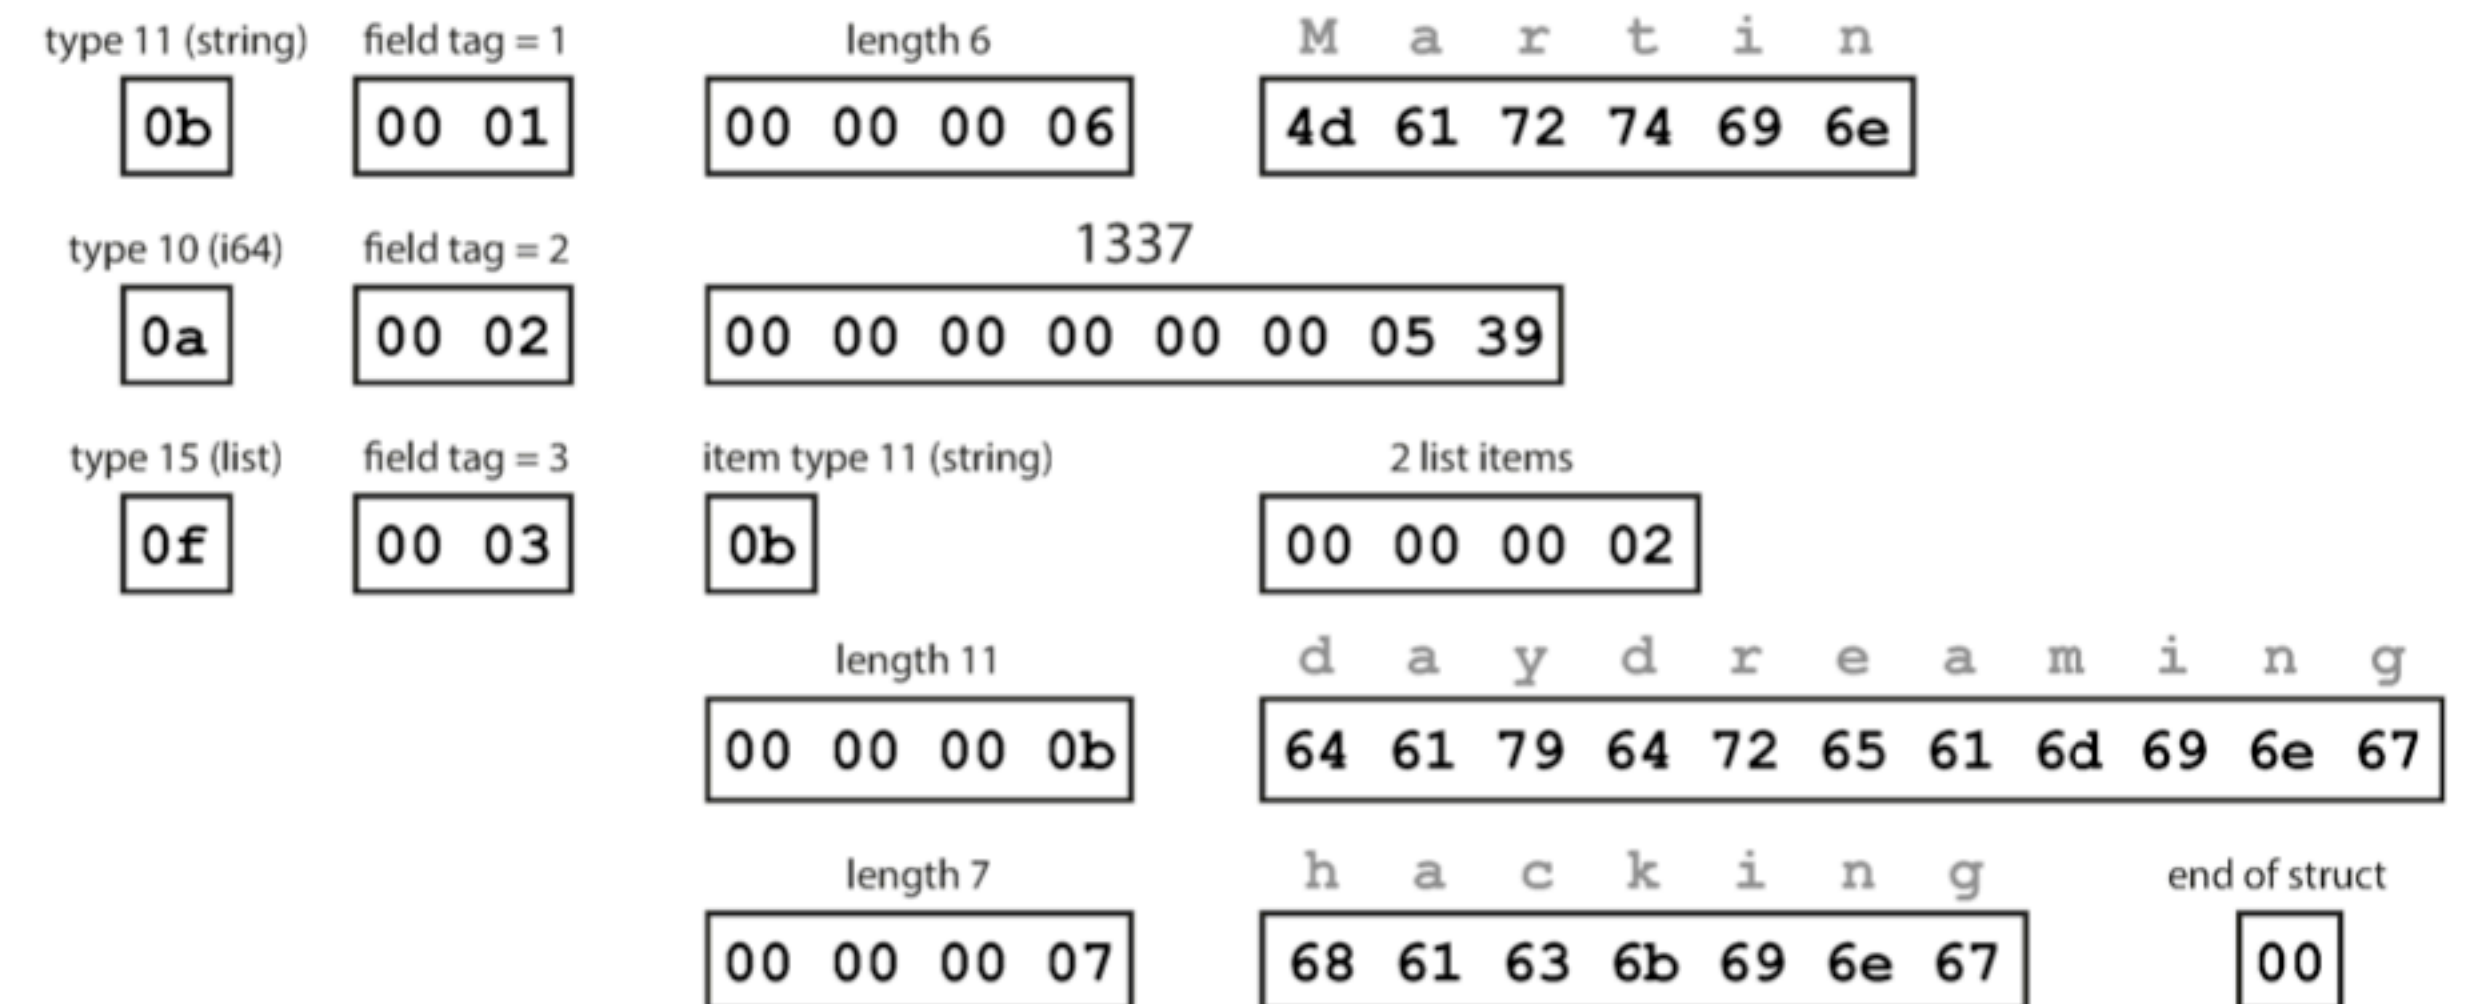
```
struct Person {  
  1: required string      userName,  
  2: optional i64         favoriteNumber,  
  3: optional list<string> interests  
}
```

Thrift BinaryProtocol

Byte sequence (59 bytes):

0b	00 01	00 00 00 06	4d 61 72 74 69 6e	0a	00 02	00 00 00 00
00 00 05 39	0f	00 03	0b	00 00 00 02	00 00 00 0b	64 61 79 64
72 65 61 6d 69 6e 67	00 00 00 07	68 61 63 6b 69 6e 67	00			

Breakdown:



Thrift Binary Protocol v.s. MessagePack

- Same:
 - each field has a type annotation
 - a length indication
 - strings also encoded as ASCII
- Diff:
 - there are no field names
 - Instead, contains field tags
 - 59 bytes vs. 81 bytes

More system performance


- Protobuf, Thrift CompactProtocol
- Key ideas:
 - Packing the field type and tag number into a single byte
 - Using variable-length integer.

Protocol Buffers

Byte sequence (33 bytes):

0a	06	4d	61	72	74	69	6e	10	b9	0a	1a	0b	64	61	79	64	72	65	61
6d	69	6e	67	1a	07	68	61	63	6b	69	6e	67							

Breakdown:



Schema evolution:

- Field tags
 - to maintain backward compatibility, every field you add after the initial deployment of the schema must be optional or have a default value.
 - only remove a field that is optional (a required field can never be removed)
 - never use the same tag number again
- Data types:
 - Possible but huge cost. May lose precision or get truncated.
 - Many language specific tricks.

The Merits of Schemas

- more compact than the “binary JSON” variants => omit field names.
- The schema & documentation.
- Hard to manually maintained.
- Keeping a database of schemas allows you to check forward and backward compatibility of schema changes.
- Code generation.

Today's topic: Data Encoding

- Formats for Encoding Data
 - Language-Specific Formats
 - JSON, XML, and Binary Variants
 - BINARY ENCODING
- Modes of dataflow
 - Database
 - REST, RPC, GraphQL
 - Message-passing
- Summary

Modes of dataflow

- Dataflow:
 - Encoding + Sharing (flowing) + Decoding
- Via databases (see “Dataflow Through Databases”)
- Via service calls (see “Dataflow Through Services: REST and RPC”)
- Via asynchronous message passing (see “Message-Passing Dataflow”)

Dataflow Through Databases

- The write process encodes data
- The read process decodes data
- Can be the same process.

